

AudioMuse-AI: Full Stack Local Deployment Guide with Docker Compose on Debian

This guide will walk you through deploying your complete AudioMuse-AI application stack on a Debian-based local machine using Docker Compose. This approach is ideal if you don't have a Kubernetes environment set up and want to run all services (Flask app, RQ worker, Redis, PostgreSQL) using pre-built Docker images, simplifying the setup process.

Prerequisites

Before you begin, ensure your Debian machine meets the following requirements:

- **Debian Operating System:** This guide is tailored for Debian 10 (Buster), Debian 11 (Bullseye), or Debian 12 (Bookworm).
- **Internet Connection:** Required to download Docker and container images.
- **Sudo Privileges:** You'll need `sudo` access to install software.

Step-by-Step Deployment

Step 1: Install Docker

First, you need to install Docker Engine on your Debian machine.

1. Update your system's package index:

```
1 sudo apt update
```

Listing 1: Update apt package index

2. Install necessary packages to allow apt to use a repository over HTTPS:

```
1 sudo apt install ca-certificates curl gnupg lsb-release -y
```

Listing 2: Install HTTPS transport packages

3. Add Docker's official GPG key:

```
1 sudo mkdir -p /etc/apt/keyrings
2 curl -fsSL https://download.docker.com/linux/debian/gpg | sudo gpg --dearmor -o /etc/apt/
  keyrings/docker.gpg
```

Listing 3: Add Docker GPG key

4. Set up the Docker repository:

```
1 echo \
2 "deb [arch=$(dpkg --print-architecture) signed-by=/etc/apt/keyrings/docker.gpg] https://
  download.docker.com/linux/debian \
3 "${lsb_release -cs} stable" | sudo tee /etc/apt/sources.list.d/docker.list > /dev/null
```

Listing 4: Configure Docker repository

5. Install Docker Engine, containerd, and Docker Compose:

```
1 sudo apt update
2 sudo apt install docker-ce docker-ce-cli containerd.io docker-buildx-plugin docker-compose
  -plugin -y
```

Listing 5: Install Docker components

6. Verify Docker installation:

```
1 sudo docker run hello-world
```

Listing 6: Verify Docker

You should see a message confirming Docker is working.

7. Add your user to the docker group (optional, but recommended to run Docker without sudo):

```
1 sudo usermod -aG docker $USER
2 newgrp docker # Apply group changes immediately, or log out and back in
```

Listing 7: Add user to docker group

Step 2: Create Project Directory and Docker Compose File

Instead of cloning the repository, you'll create a local directory and the `docker-compose.yaml` file within it.

1. Create a directory for your project:

```
1 mkdir ~/audiomuse-ai
2 cd ~/audiomuse-ai
```

Listing 8: Create project directory

2. Create the Docker Compose file (docker-compose.yaml): Create a file named `docker-compose.yaml` (no extension) in your `audiomuse-ai` project directory and paste the following content into it. This file defines all the services needed for your application, using the pre-built Docker image.

```
1 version: '3.8'
2
3 services:
4   # Redis service for RQ (task queue)
5   redis:
6     image: redis:7-alpine
7     container_name: audiomuse-redis
8     ports:
9       - "6379:6379" # Expose Redis port to the host
10    volumes:
11      - redis-data:/data # Persistent storage for Redis data
12    restart: unless-stopped
13
14   # PostgreSQL database service
15   postgres:
16     image: postgres:15-alpine
17     container_name: audiomuse-postgres
18     environment:
19       POSTGRES_USER: "audiomuse"
20       POSTGRES_PASSWORD: "audiomusepassword"
21       POSTGRES_DB: "audiomusedb"
22     ports:
23       - "5432:5432" # Expose PostgreSQL port to the host
24     volumes:
```

```

25     - postgres-data:/var/lib/postgresql/data # Persistent storage for PostgreSQL data
26     restart: unless-stopped
27
28 # AudioMuse-AI Flask application service
29 audiomuse-ai-flask:
30     image: ghcr.io/neptunehub/audiomuse-ai:0.2.0-alpha # Your pre-built image
31     container_name: audiomuse-ai-flask-app
32     ports:
33         - "8000:8000" # Map host port 8000 to container port 8000
34     environment:
35         SERVICE_TYPE: "flask" # Tells the container to run the Flask app
36         JELLYFIN_USER_ID: "0e45c44b3e2e4da7a2be11a72a1c8575" # From jellyfin-credentials
37             secret
38         JELLYFIN_TOKEN: "e0b8c325bc1b426c81922b90c0aa2ff1" # From jellyfin-credentials secret
39         JELLYFIN_URL: "http://jellyfin.192.168.3.131.nip.io:8087" # From audiomuse-ai-config
40             ConfigMap
41         DATABASE_URL: "postgresql://audiomuse:audiomusepassword@postgres:5432/audiomusedb" #
42             Connects to the 'postgres' service
43         REDIS_URL: "redis://redis:6379/0" # Connects to the 'redis' service
44         TEMP_DIR: " temp_audio"
45     volumes:
46         - temp-audio-flask: temp_audio # Volume for temporary audio files
47     depends_on:
48         - redis
49         - postgres
50     restart: unless-stopped
51
52 # AudioMuse-AI RQ Worker service
53 audiomuse-ai-worker:
54     image: ghcr.io/neptunehub/audiomuse-ai:0.2.0-alpha # Your pre-built image
55     container_name: audiomuse-ai-worker-instance
56     environment:
57         SERVICE_TYPE: "worker" # Tells the container to run the RQ worker
58         JELLYFIN_USER_ID: "0e45c44b3e2e4da7a2be11a72a1c8575" # From jellyfin-credentials
59             secret
60         JELLYFIN_TOKEN: "e0b8c325bc1b426c81922b90c0aa2ff1" # From jellyfin-credentials secret
61         JELLYFIN_URL: "http://jellyfin.192.168.3.131.nip.io:8087" # From audiomuse-ai-config
62             ConfigMap
63         DATABASE_URL: "postgresql://audiomuse:audiomusepassword@postgres:5432/audiomusedb" #
64             Connects to the 'postgres' service
65         REDIS_URL: "redis://redis:6379/0" # Connects to the 'redis' service
66         TEMP_DIR: " temp_audio"
67     volumes:
68         - temp-audio-worker: temp_audio # Volume for temporary audio files
69     depends_on:
70         - redis
71         - postgres
72     restart: unless-stopped
73
74 # Define volumes for persistent data and temporary files
75 volumes:
76     redis-data:
77     postgres-data:
78     temp-audio-flask: # Volume for Flask app's temporary audio
79     temp-audio-worker: # Volume for Worker's temporary audio

```

Listing 9: docker-compose.yaml

Important Notes for your docker-compose.yaml:

- **Image Source:** The image directive now points directly to `ghcr.io/neptunehub/audiomuse-ai:0.2.0-alpha` meaning Docker will pull this pre-built image instead of building it locally.
- **Environment Variables:** I've directly included the values from your Kubernetes `Secret` and

`ConfigMap` into the `environment` section. For a production environment, you might prefer using a `.env` file for sensitive credentials.

- **Service Communication:** Services within a Docker Compose network can communicate using their service names (e.g., `postgres` for the PostgreSQL service, `redis` for the Redis service). This is why `DATABASE_URL` and `REDIS_URL` refer to `postgres:5432` and `redis:6379` respectively.
- **command Removed:** The `command` override for `audiomuse-ai-flask` and `audiomuse-ai-worker` has been removed. The `CMD` instruction within the pre-built Docker image (which uses the `SERVICE_TYPE` environment variable) will now correctly determine whether to run `app.py` or `rq_worker.py`.
- **Volumes:** Persistent volumes (`redis-data`, `postgres-data`) are defined to ensure your Redis data and PostgreSQL data are not lost when containers are stopped or removed. `emptyDir` equivalents (`temp-audio-flask`, `temp-audio-worker`) are used for temporary audio storage.
- **depends_on:** This ensures that `redis` and `postgres` services are started before the Flask app and RQ worker, as they depend on these services.

Step 3: Run the Docker Compose Stack

Once your `docker-compose.yaml` is in place, you can pull the images and run all services with a single command. Make sure you are in the directory where you created `docker-compose.yaml` (`~/audiomuse-ai`).

```
1 docker compose up -d
```

Listing 10: Run Docker Compose stack

- `docker compose up`: The command to start and run the services defined in `docker-compose.yaml`.
- `-d`: Runs the containers in "detached" mode (in the background).

This process will download the necessary Docker images (Redis, PostgreSQL, and your AudioMuse-AI image) and then start all the services. It might take a few minutes, depending on your internet speed.

Step 4: Access AudioMuse-AI

Your AudioMuse-AI application services should now be running in Docker containers.

(a) **Check container status:**

```
1 docker compose ps
```

Listing 11: Check container status

You should see `audiomuse-redis`, `audiomuse-postgres`, `audiomuse-ai-flask-app`, and `audiomuse-ai-worker` listed with status `running`.

(b) **View logs (optional, for debugging):** To see the combined logs of all services:

```
1 docker compose logs -f
```

Listing 12: View combined logs

To see logs for a specific service (e.g., the Flask app):

```
1 docker compose logs -f audiomuse-ai-flask
```

Listing 13: View specific service logs

(c) **Access the application:** Open your web browser and navigate to:

`http://localhost:8000`

Managing Your Docker Compose Stack

Here are some useful Docker Compose commands for managing your application stack:

- **Stop all services (without removing containers):**

```
1 docker compose stop
```

Listing 14: Stop all services

- **Start all services (if stopped):**

```
1 docker compose start
```

Listing 15: Start all services

- **Stop and remove all services, networks, and volumes (clean up):**

```
1 docker compose down -v
```

Listing 16: Stop and remove all services

- `-v`: Removes volumes as well. Use this if you want a fresh start and don't care about persistent data (Redis, PostgreSQL).

- **Restart all services:**

```
1 docker compose restart
```

Listing 17: Restart all services

- **Pull updated images and restart services (if a new version of your image is released):**

```
1 docker compose pull && docker compose up -d
```

Listing 18: Pull and restart services

- **Scaling the RQ Worker (e.g., to run 2 workers):** To ensure proper functionality, your AudioMuse-AI application requires a minimum of **two** RQ worker instances. One worker is typically dedicated to handling the main tasks, while another handles any subtasks or child processes that might be spawned. Running with fewer than two workers might lead to tasks getting stuck or the application not functioning as expected.

To run two (or more) instances of the `audiomuse-ai-worker` service, use the `--scale` flag with `docker compose up`. This will create additional containers for the specified service.

First, ensure your existing services are stopped or down:

```
1 docker compose down
```

Listing 19: Stop existing services

Then, start all services and scale the worker:

```
1 docker compose up -d --scale audiomuse-ai-worker=2
```

Listing 20: Scale worker service

You can replace 2 with any higher number of worker instances if you need more processing power. To verify, run `docker compose ps` and you should see multiple `audiomuse-ai-worker-instance` containers.

This updated guide provides a comprehensive and robust way to deploy your full AudioMuse-AI application locally using Docker Compose, directly pulling the pre-built image and simplifying the setup process significantly.